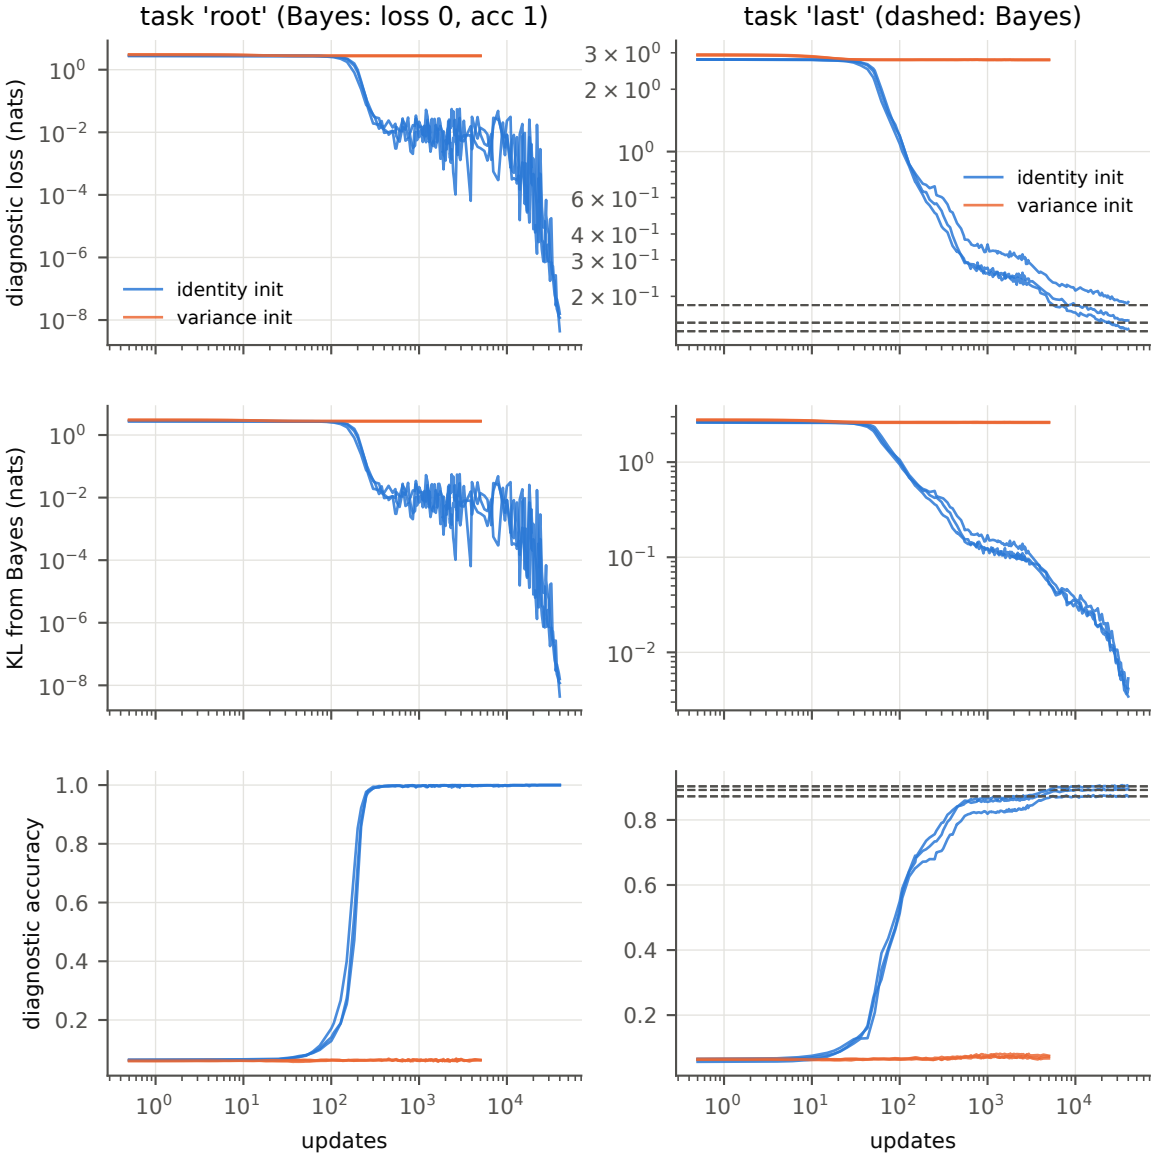
Identity vs variance init, r = 16, m = 4, init seed 0 (one line per grammar)
task 'root' (Bayes: loss 0, acc 1)
task 'last' (dashed: Bayes)
diagnostic loss (nats)
KL from Bayes (nats)
diagnostic accuracy
identity init
variance init
updates