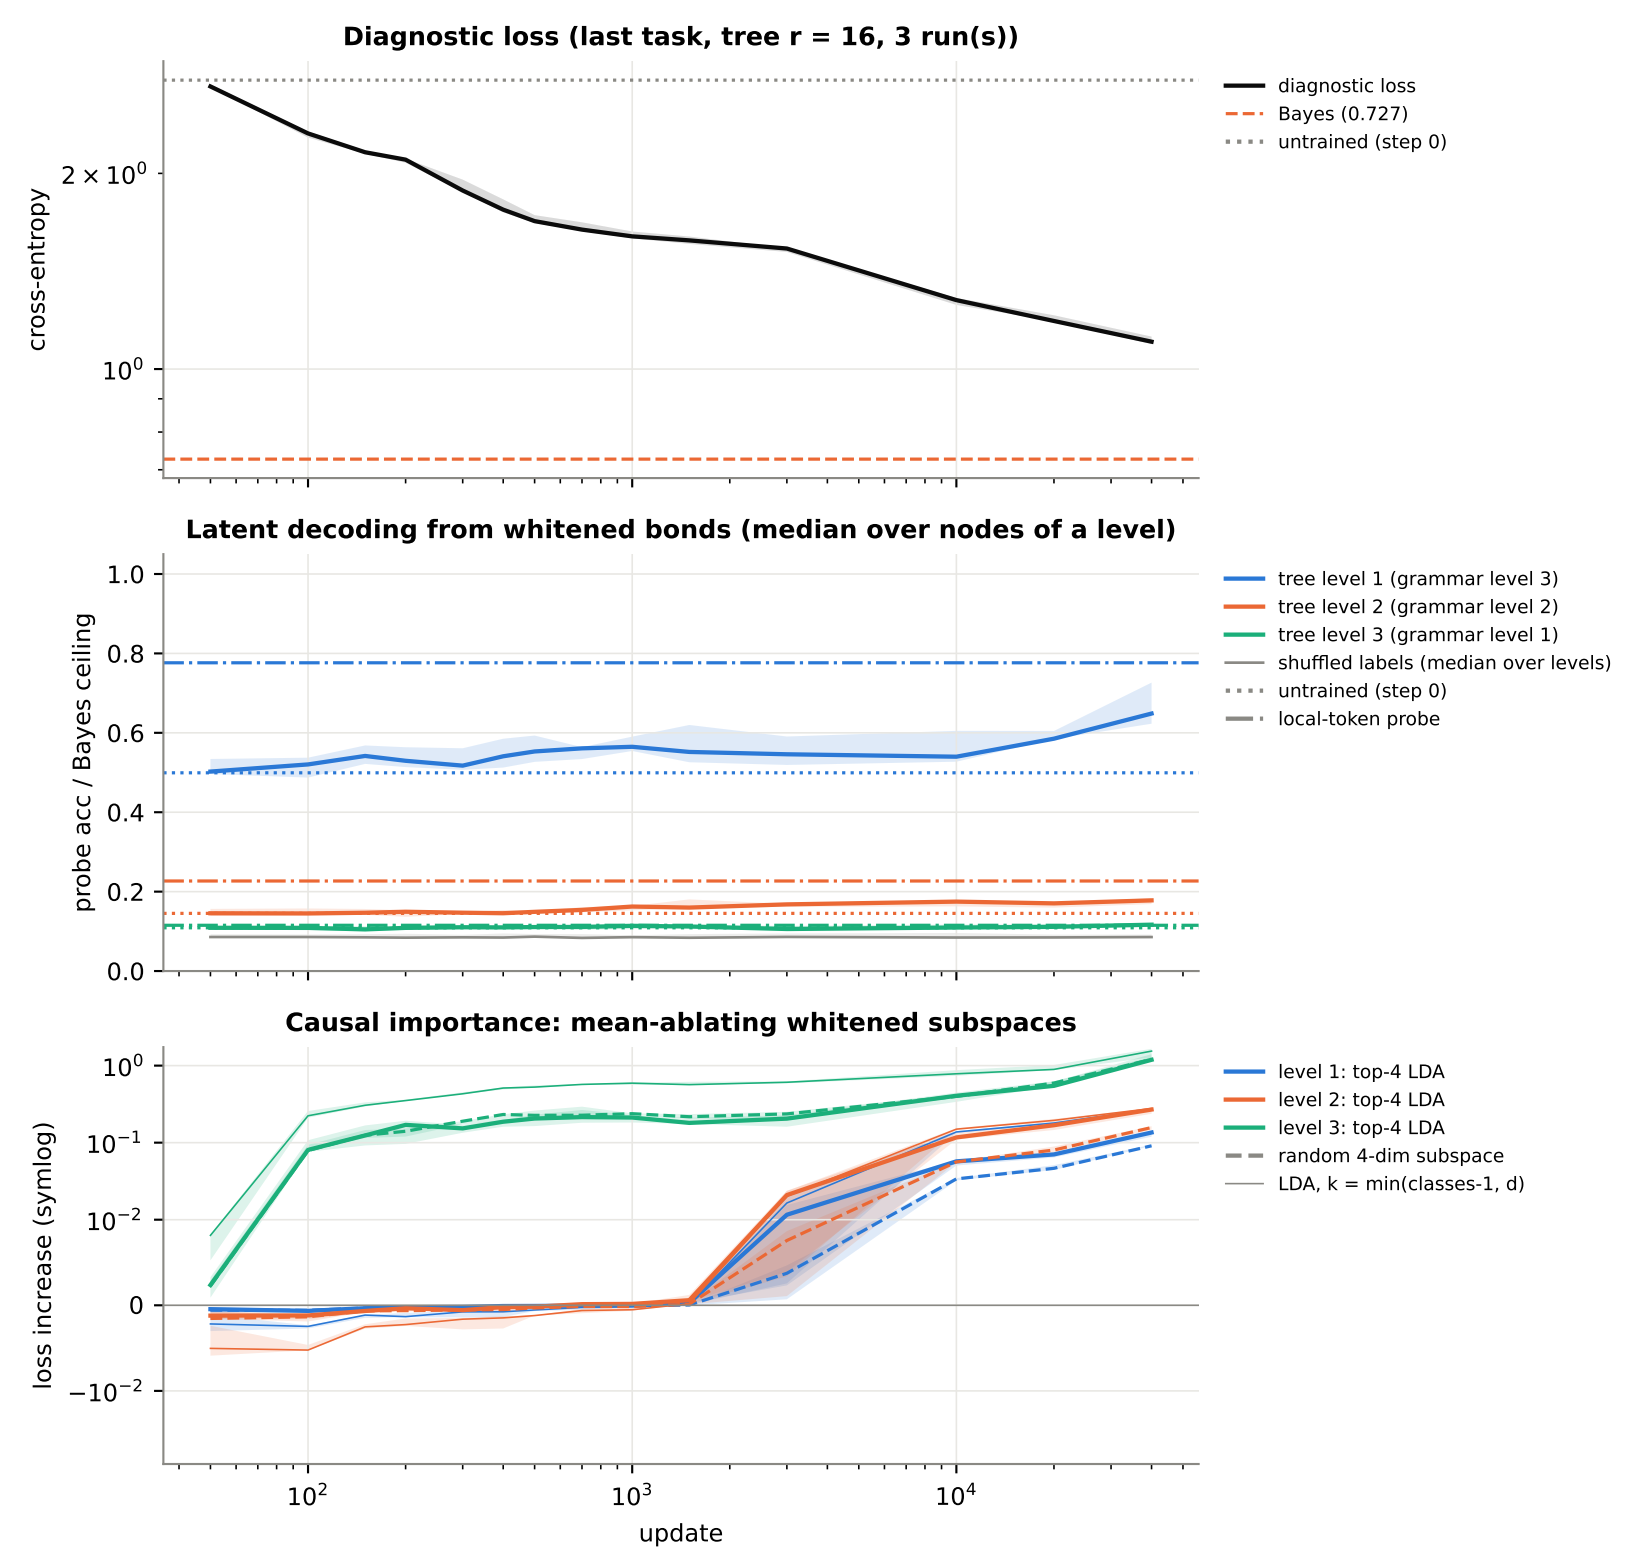
Diagnostic loss (last task, tree r = 16, 3 run(s))
cross-entropy
2×10^0
10^0
diagnostic loss
Bayes (0.727)
untrained (step 0)
Latent decoding from whitened bonds (median over nodes of a level)
probe acc / Bayes ceiling
1.0
0.8
0.6
0.4
0.2
0.0
tree level 1 (grammar level 3)
tree level 2 (grammar level 2)
tree level 3 (grammar level 1)
shuffled labels (median over levels)
untrained (step 0)
local-token probe
Causal importance: mean-ablating whitened subspaces
loss increase (symlog)
10^0
10^{-1}
10^{-2}
0
-10^{-2}
10^2
10^3
10^4
update
level 1: top-4 LDA
level 2: top-4 LDA
level 3: top-4 LDA
random 4-dim subspace
LDA, k = min(classes-1, d)