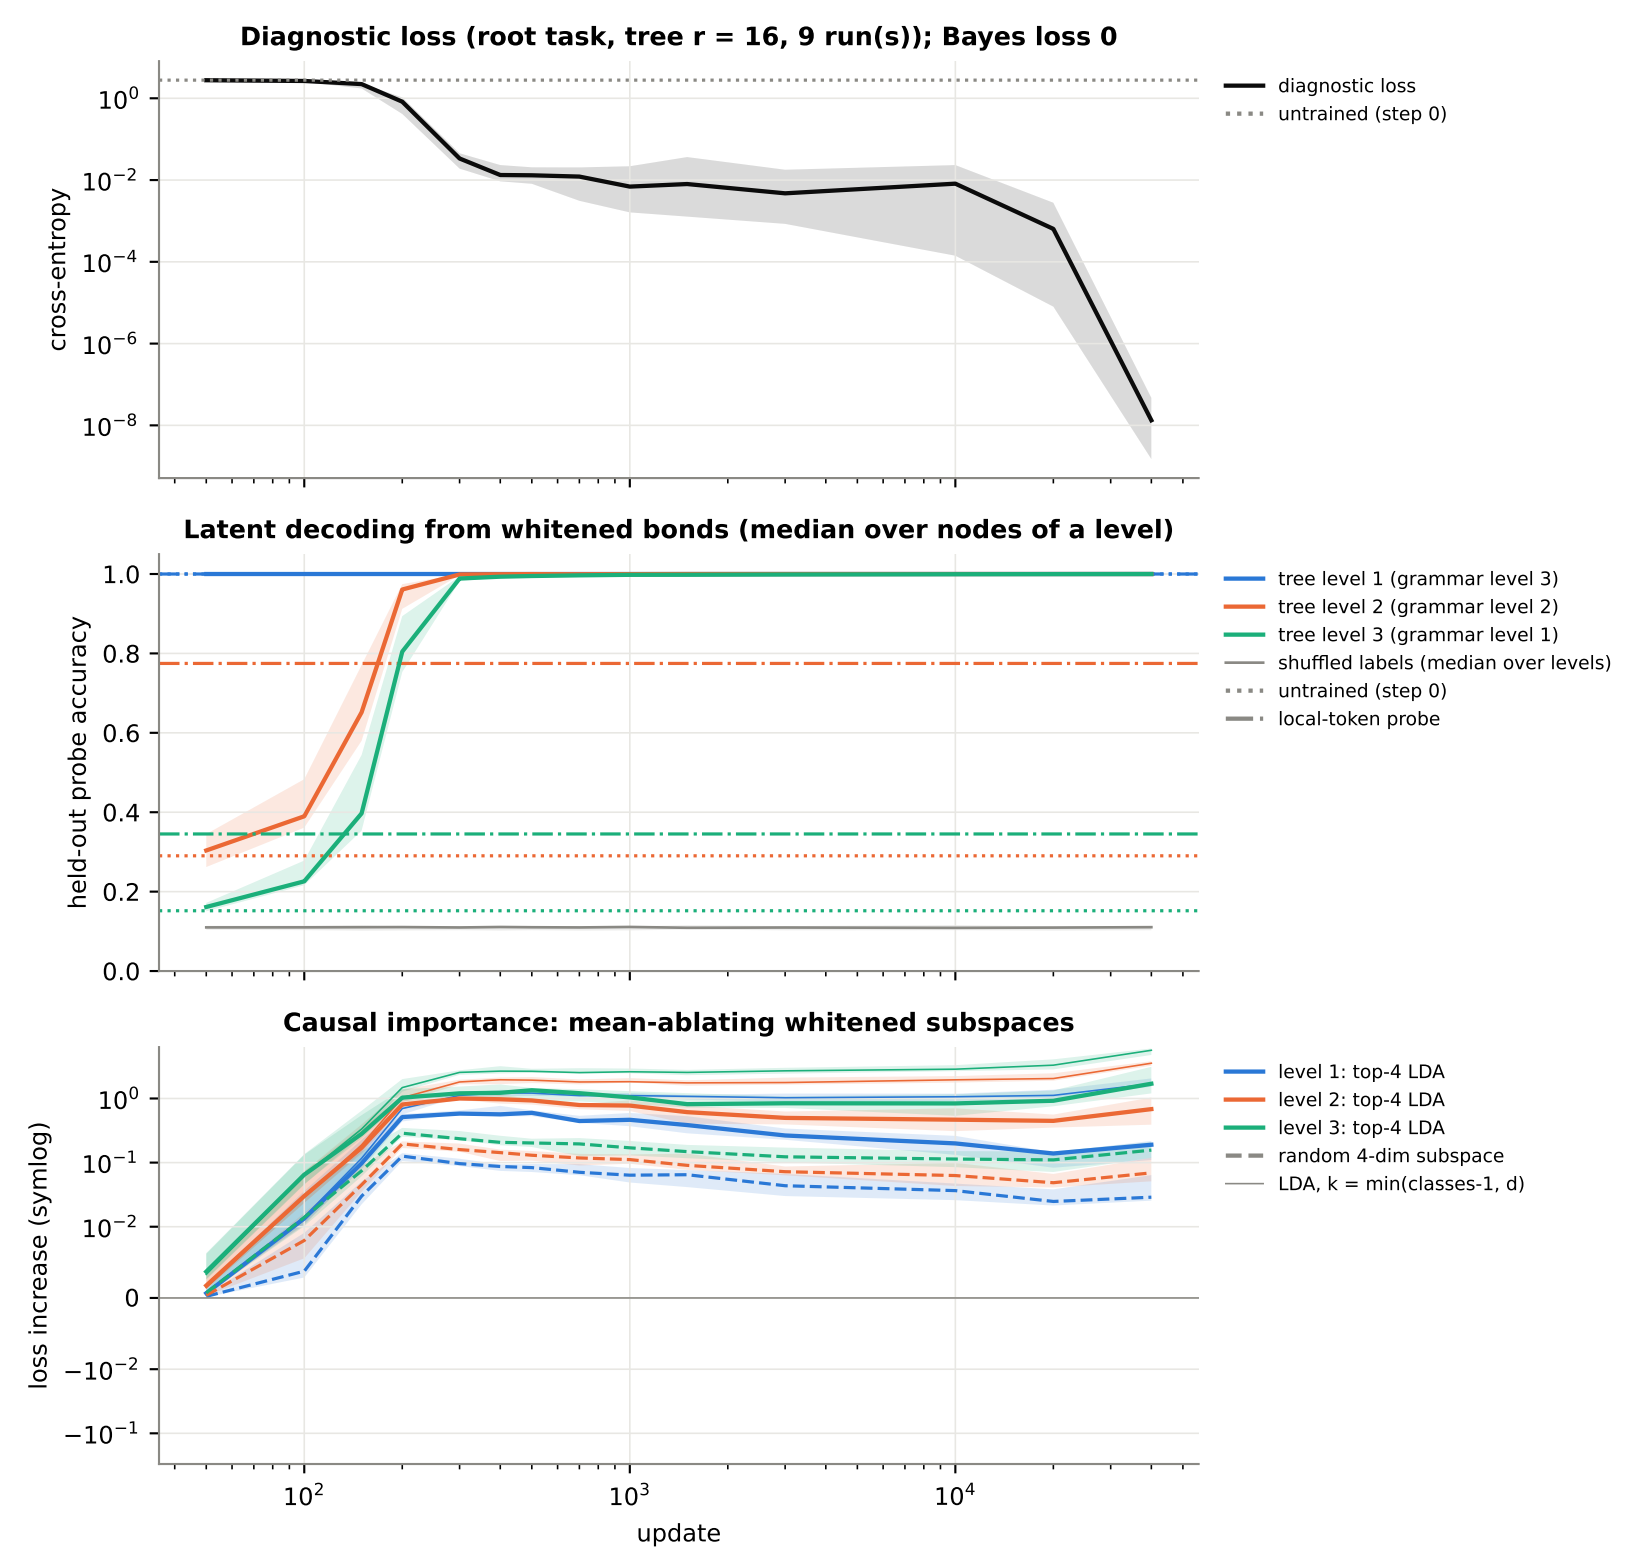
Diagnostic loss (root task, tree r = 16, 9 run(s)); Bayes loss 0
cross-entropy
diagnostic loss
untrained (step 0)
Latent decoding from whitened bonds (median over nodes of a level)
held-out probe accuracy
tree level 1 (grammar level 3)
tree level 2 (grammar level 2)
tree level 3 (grammar level 1)
shuffled labels (median over levels)
untrained (step 0)
local-token probe
Causal importance: mean-ablating whitened subspaces
loss increase (symlog)
level 1: top-4 LDA
level 2: top-4 LDA
level 3: top-4 LDA
random 4-dim subspace
LDA, k = min(classes-1, d)
update